An R toolbox for score-based measurement invariance tests in IRT models

Lennart Schneider[1,2] · Carolin Strobl[3] · Achim Zeileis[4] · Rudolf Debelak[3,5]

Abstract

The detection of differential item functioning (DIF) is a central topic in psychometrics and educational measurement. In the past few years, a new family of score-based tests of measurement invariance has been proposed, which allows the detection of DIF along arbitrary person covariates in a variety of item response theory (IRT) models. This paper illustrates the application of these tests within the R system for statistical computing, making them accessible to a broad range of users. This presentation also includes IRT models for which these tests have not previously been investigated, such as the generalized partial credit model. The paper has three goals: First, we review the ideas behind score-based tests of measurement invariance. Second, we describe the implementation of these tests within the R system for statistical computing, which is based on the interaction of the R packages *mirt*, *psychotools* and *strucchange*. Third, we illustrate the application of this software and the interpretation of its output in two empirical datasets. The complete R code for reproducing our results is reported in the paper.

Keywords Differential item functioning · Item response theory · Software tutorial

Introduction

A property of psychological and educational tests that is desirable in many applications is that test takers of equal ability show the same probability of giving correct or positive responses to the individual items, and that this probability does not depend on other variables. This property is typically termed measurement invariance (Millsap, 2011). On a conceptual level, an absence of measurement invariance can indicate that a test measures irrelevant characteristics of the respondents and that its fairness is violated (Camilli, 2006). Many authors and international testing standards thus emphasize the importance of checking the presence of differential item functioning (DIF; Holland & Wainer, 1993) in psychological measurement (American Educational Research Association, 2014; Borsboom, 2006). We illustrate this problem by a typical example, which we will revisit later on in the paper:

Consider a questionnaire for the assessment of verbal aggression, as presented in the Verbal Aggression dataset that was described in the book of De Boeck and Wilson (2004) and the paper of De Boeck et al. (2011). An example item is "A bus fails to stop for me. I would want to curse.". Item response theory (IRT) models can be used to describe the interaction of test takers and test items in such a scenario (De Boeck et al., 2004, 2011). Here, DIF corresponds to a situation where respondents from various subgroups that have the same level of verbal aggression differ with regard to their tendency to agree to items from this test. Often these subgroups of respondents can be described by combinations of person characteristics, which are termed covariates in the following.

Over the last decades, various tests for the detection of DIF have been proposed. Among the suggested methods, several families of tests can be discerned. Traditional approaches are based on the comparison of predefined groups of respondents (for an overview, see Magis et al.,

✉ Rudolf Debelak
  rudolf.debelak@psychologie.uzh.ch

[1] University of Tübingen, Tübingen, Germany

[2] Present address: Ludwig Maximilian University of Munich, Munich, Germany

[3] Department of Psychology, University of Zurich, Zurich, Switzerland

[4] Universität Innsbruck, Innsbruck, Austria

[5] Institute of Psychology, University of Leipzig, Neumarkt 9, 04109 Leipzig, Germany

2010). For this purpose, a reference group and one or more focal groups are defined before the analysis. A practical disadvantage of these traditional approaches is that it remains unclear how to address the problem of detecting DIF effects if no natural groups are available. This can occur, for instance, if the respondents differ with regard to continuous covariates such as age.

To address this problem, several alternative approaches were proposed to assess measurement invariance with respect to non-categorical person covariates. Liu et al. (2016) proposed a semiparametric IRT model and a permutation test which allows the detection of DIF effects with regard to a continuous covariate. Related approaches proposed IRT models which included person covariates in the item response functions (e.g., Moustaki, 2003). These models either require estimation of the form of the item response curve, or assume a specific relationship between the person covariate and the observed responses.

A second alternative approach is based on moderated factor analysis (Bauer, 2017, 2009, Molenaar, 2020). Here, person covariates are included in a factor model, which allows the modeling of changes of factor loadings and other model parameter with regard to person covariates. Moderation functions are used to describe the relationship between the observed person covariates and the other model parameters.

A third alternative approach is based on mixture distribution models. An important example is the mixed Rasch model (Rost, 1990). Here, a number of latent groups of respondents is assumed. Respondents from different groups differ with regard to the specific values of the item parameters. Whereas traditional mixture distribution models do not define a relationship between these latent groups and observed covariates, several authors have suggested extensions of mixture distribution models which include person covariates (Dai, 2013; Li et al., 2016; Tay et al., 2011). The resulting models allow an estimation of the relationship between observed covariates and the probability of individual respondents for being a member in a specific latent group.

In this paper, we present a family of score-based measurement invariance tests that can be used to check the presence of DIF with regard to observed continuous, ordinal and categorical covariates, such as age, educational level or gender. The resulting family of tests is very flexible and does not assume a specific relationship between the observed person covariates and possible DIF effects. These tests can be applied to a wide range of IRT models, both with conditional maximum likelihood (CML) and marginal maximum likelihood (MML) estimation methods (Baker & Kim, 2004). Table 1 gives an overview of the supported IRT models. Note that in this paper, we use IRT as a broad umbrella term for item response models that also includes Rasch-type models.

As can be seen from Table 1, the software presented in this paper allows the calculation of score-based tests for various IRT models for which these tests have not been described or derived before. In summary, this tutorial paper provides the following contributions: First, we review the ideas behind score-based tests of measurement invariance. Second, we present their implementation in the R system for statistical computing (R Core Team, 2019). Third,

Table 1 Overview of supported IRT models

Model	Response type	Estimation	Model reference	Score-based tests published in
Rasch Model	dichotomous	CML	Rasch (1960)	Strobl et al. (2015)
1PL	dichotomous	MML	Birnbaum (1968)	
2PL	dichotomous	MML	Birnbaum (1968)	Debelak and Strobl (2019)
3PL	dichotomous	MML	Birnbaum (1968)	Debelak and Strobl (2019)
3PLu	dichotomous	MML	Barton and Lord (1981)	
4PL	dichotomous	MML	Barton and Lord (1981)	
ideal point model	dichotomous	MML	Maydeu-Olivares et al. (2006)	
rating scale model	polytomous	CML	Andrich (1978)	Komboz et al. (2018)
partial credit model	polytomous	CML	Masters (1982)	Komboz et al. (2018)
generalized partial credit model	polytomous	MML	Muraki (1992)	
graded response model	polytomous	MML	Samejima (1969)	
nominal response model	polytomous	MML	Bock (1972)	
generalized graded unfolding model	polytomous	MML	Roberts et al. (2000)	
monotonic polynomial model	polytomous	MML	Falk and Cai (2016)	

Note. For all models that can be fitted via MML using the R package *mirt* (Chalmers, 2012), mixed itemtype models as well as multidimensional confirmatory and exploratory versions and multiple group versions of the models are supported. Moreover, arbitrary parameter constraints are possible. This allows for fitting, e.g., a partial credit model via MML

we illustrate the application of this software and the interpretation of its output by means of two empirical datasets. The audience of this paper includes substantive researchers who wish to apply these methods to their own datasets as well as psychometricians who are interested in their theoretical background.

The following section "A Conceptual and Formal Framework for Score-Based Measurement Invariance Tests" briefly describes the underlying framework of these tests. The section "The Implementation of Score-Based Measurement Invariance Tests within R" describes their implementation within R (R Core Team, 2019). In the section "Illustrations with Empirical Data", we illustrate this methodology by means of two datasets in a tutorial style. This section includes an analysis of the Verbal Aggression dataset described by De Boeck and Wilson (2004) and De Boeck et al. (2011) as well as an analysis of the Generic Conspiracist Beliefs dataset (Brotherton et al., 2013; Open Source Psychometrics Project, 2016). We conclude our paper with a general discussion in the final section.

A conceptual and formal framework for score-based measurement invariance tests

An overview

The class of score-based tests considered in the following was originally developed in the econometrics literature as a method to test for instability (or "fluctuation") in ordinary least squares (OLS) or maximum likelihood (ML) estimates (Andrews, 1993). They can also be applied to other types of M-estimators (beyond OLS and ML), which is why they are named M-fluctuation tests in some publications (Strobl et al., 2015; Zeileis & Hornik, 2007). In the field of psychometrics, they were first described for Bradley-Terry models (Strobl et al., 2011), factor analytical models (Merkle et al., (Merkle & Zeileis, 2013; Merkle et al., 2014)) and the dichotomous Rasch model (Strobl et al., 2015). Later works generalized this method to other models of item response theory (Debelak & Strobl, 2019; Komboz et al., 2018; Wang et al., 2018), including the application of these tests in a model-based recursive partitioning framework to detect subpopulations that show measurement invariance (Komboz et al., 2018; Strobl et al., 2015).

The family of score-based tests can be considered a generalization of the classical Lagrange multiplier test that was applied by Glas and colleagues to detect DIF effects in various IRT models when pre-defined groups are available (Glas, 1998, 1999, 2003). Another closely related approach for checking the invariance of individual model parameters in dynamic panel models that is based on regression analysis

was recently proposed by Arnold et al. (2020), using ideas presented by Oberski (2013).

To present the formal ideas behind these tests, we start by considering an IRT model whose model parameters are summarized by a vector $\boldsymbol{\Psi} = (\Psi_1, \ldots, \Psi_P)$ and a matrix of observed responses U. We further denote the rows of U, which correspond to responses given by N individual respondents, by $(\boldsymbol{u}_1, \ldots, \boldsymbol{u}_N)$. It follows from the local independence assumption made in most IRT models that the log-likelihood $\ell(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N)$ of these models is a sum of case-wise contributions:

$$\ell(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N) = \sum_{i=1}^{N} \ell(\boldsymbol{\Psi}; \boldsymbol{u}_i).$$

Score-based tests use the scores of the model parameters to assess their invariance. Here the score is the gradient, i.e., the vector of first partial derivatives of the log-likelihood with respect to a model parameter of interest. Using our notation, the score is thus given by:

$$s(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N) = \left(\frac{\partial \ell(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N)}{\partial \Psi_1}, \ldots, \frac{\partial \ell(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N)}{\partial \Psi_P} \right)^t.$$

As was the case for the log-likelihood, we can present the score as a sum of independent, individual score contributions:

$$s(\boldsymbol{\Psi}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N) = \sum_{i=1}^{N} s(\boldsymbol{\Psi}; \boldsymbol{u}_i),$$

where

$$s(\boldsymbol{\Psi}; \boldsymbol{u}_i) = \left(\frac{\partial \ell(\boldsymbol{\Psi}; \boldsymbol{u}_i)}{\partial \Psi_1}, \ldots, \frac{\partial \ell(\boldsymbol{\Psi}; \boldsymbol{u}_i)}{\partial \Psi_P} \right)^t.$$

If we replace $\boldsymbol{\Psi}$ by its maximum likelihood estimate $\hat{\boldsymbol{\Psi}}$, we directly obtain:

$$s(\hat{\boldsymbol{\Psi}}; \boldsymbol{u}_1, \ldots, \boldsymbol{u}_N) = \sum_{i=1}^{N} s(\hat{\boldsymbol{\Psi}}; \boldsymbol{u}_i) = \boldsymbol{0}.$$

As can be seen, the sum over all individual score contributions, evaluated at the maximum likelihood estimate, is $\boldsymbol{0}$, that is, a vector of zeros. The reason is that by definition the score is $\boldsymbol{0}$ at the maximum likelihood estimator, where the log-likelihood reaches its maximum. Under the null hypothesis that the model underlying the estimation is the true model and all parameters are stable, we can also determine the asymptotic distribution of cumulative sums of scores. If the parameters are not invariant, the sum of the individual score contributions is still $\boldsymbol{0}$, but unexpected deviations from this asymptotic distribution occur. The basic idea of the score-based invariance tests is to first sort all respondents with regard to a person covariate of interest. We denote the

resulting ordered responses by $(\boldsymbol{u}_{(1)}, \ldots, \boldsymbol{u}_{(N)})$. Now, we consider the following term for a value t between 0 an 1:

$$B(t, \hat{\boldsymbol{\Psi}}) = \hat{I}^{-1/2} N^{-1/2} \sum_{i=1}^{\lfloor Nt \rfloor} s(\hat{\boldsymbol{\Psi}}; \boldsymbol{u}_{(i)}).$$

Here, $\lfloor \rfloor$ is the floor function and $\hat{I}$ denotes a consistent estimate of the covariance matrix of the case-wise score contributions. Its inclusion leads to a standardization of $B(t, \hat{\boldsymbol{\Psi}})$.

If no DIF is present, the expected distribution of $s(\hat{\boldsymbol{\Psi}}; \boldsymbol{u}_{(i)})$ will be independent from the index (i), that is, the ordering with respect to the person covariate. It follows that the distribution of the cumulative sum $B(t, \hat{\boldsymbol{\Psi}})$ is known for sufficiently large samples if the IRT model is correctly specified and under the null hypothesis that no DIF effects are present. Mathematically, this follows from the functional central limit theorem, or Donsker's theorem (Billingsley, 1995).

As an illustration, Fig. 1 gives two examples of the cumulative sum of scores for a single parameter in an IRT context when DIF is absent (top) or present (bottom). As can be seen, the cumulative sum fluctuates randomly around 0 if DIF is absent, but leads to a systematic pattern if DIF is present in a certain covariate. Although this illustration uses a continuous covariate, the same principle can be used for categorical and ordinal covariates.

The information on the invariance of the various model parameters can be summarized by test statistics, and the *p*-value under the null hypothesis of parameter invariance can be calculated. For computing the test statistic, we can use all available model parameters, focus on groups of model parameters or can even use only single model parameters. Conceptually, this allows the assessment of measurement invariance for the whole item set, groups of items or single items. In this tutorial, we will focus on invariance tests for the complete item set, but will also give an example for a test where invariance is checked for a single item.

Applications of score-based tests in IRT models share the following common steps for assessing the invariance of model parameters:

1. First, an appropriate model is chosen for describing the data and its parameters are estimated using a ML approach. In IRT models, this is typically done via CML or MML estimation. The user can be supported in this step by wrapper functions (see below).
2. Second, the case-wise score contributions are calculated and used for measuring the invariance of the parameter estimates with regard to some person covariate.
3. Third, this item-wise invariance information is summarized using a test statistic, either for the complete item set or a subset. The choice of test statistic typically
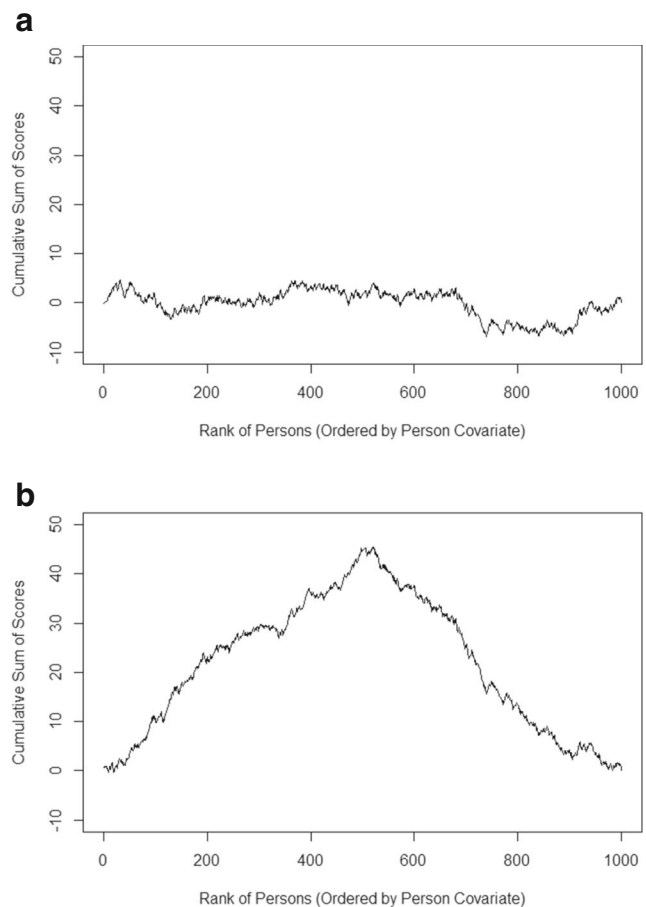
Fig. 1 **a** The cumulative sum of scores if DIF is absent. The observed values fluctuate randomly around 0. **b** The cumulative sum of scores if DIF is present between the first and the second half of the sample. Instead of a random fluctuation, a typical pattern emerges

depends on whether the person covariate that is used for the analysis is categorical, ordinal, or continuous.

For IRT models, this leads to the workflow illustrated in Fig. 2.

In the following subsections, we will discuss these steps in more detail.

Estimating the Item parameters of IRT models

Historically, numerous methods have been proposed for estimating IRT model parameters, but we focus here on two methods that are most widely used today: CML and MML estimation. CML estimation makes use of specific characteristics of Rasch-type models and can generally not be used to estimate the item parameters of more general IRT models, like the two-parametric logistic or 2PL model (Birnbaum, 1968). MML estimation can be used to estimate the item parameters in more general models, but makes specific assumptions on the person parameter distribution. For more detailed presentations, see Baker and Kim (2004)
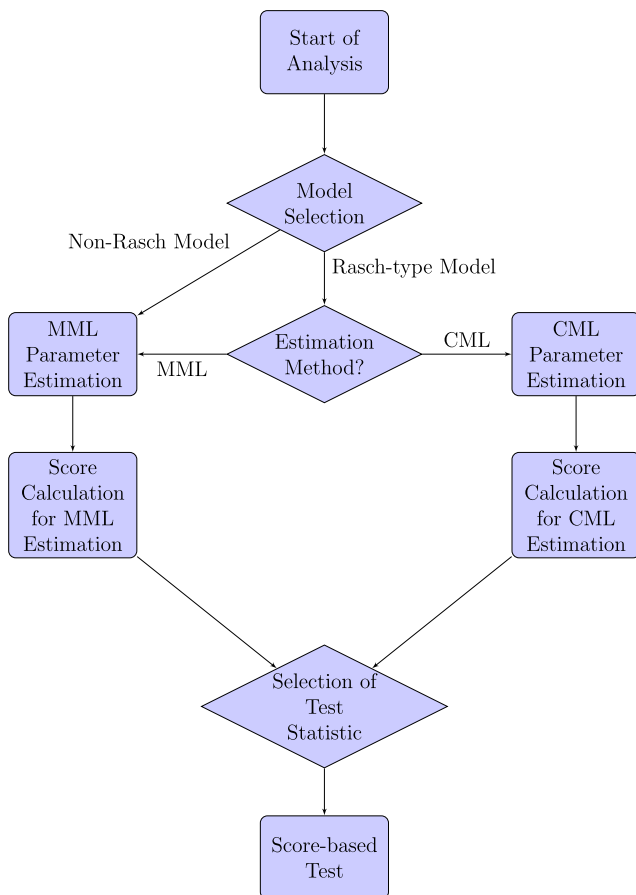
Fig. 2 The workflow of the application of score-based measurement invariance tests for IRT models. Diamond-shaped nodes correspond to decisions made by the user, whereas rectangle-shaped nodes stand for steps in the resulting type of analysis

or Molenaar ([1995](#)). Once the item parameters have been estimated, the next step of the data analysis consists of the calculation of the individual score contributions.

Determining the distributions of the individual score contributions

In IRT models, scores for individual item parameters can be described as sums over the respondent sample, with each summand being an individual score contribution. For an illustration, we consider the 2PL model in slope/intercept parametrization. Let u_{ij} denote the response of respondent i ($i = 1, \ldots, N$) to item j ($j = 1, \ldots, J$). In this model, the probability of a positive response of respondent i to item j depends on an item-specific slope parameter a_j, an item-specific intercept parameter d_j, and a respondent-specific ability parameter θ_i through the item response function:

$$P_{ij} := P(u_{ij} = 1 | a_j, d_j, \theta_i) = \frac{\exp(a_j\theta_i + d_j)}{1 + \exp(a_j\theta_i + d_j)}.$$

As an example, the derivative of the marginal log-likelihood with regard to the slope parameter takes on the following form (Baker & Kim, [2004](#); Debelak & Strobl, [2019](#)):

$$\sum_{i=1}^{N} \int (u_{ij} - P_{ij}) \cdot \theta_i \cdot P(\theta_i | \boldsymbol{u}_i, \boldsymbol{\Psi}) d\theta_i.$$

In this equation, $P(\theta_i | \boldsymbol{u}_i, \boldsymbol{\Psi})$ depends on the model parameters, which we summarize as $\boldsymbol{\Psi}$, and the responses given by respondent i, which we summarize as $\boldsymbol{u}_i$. This equation shows again that the derivative can be presented as a sum of individual contributions of the N respondents. This observation also holds for the other item parameters. It is thus possible to calculate the individual score contributions for all respondents and all item parameters given the observed response data and the item parameter estimates.

Testing measurement invariance with score-based tests

The presence of DIF leads to a systematic deviation of the observed distribution of the cumulative individual score contributions from the distribution that is expected when no DIF is present. It is therefore necessary to summarize the deviations found in all item parameters. Several test statistics have been proposed in the literature for this purpose. Here we provide a brief overview of some selected test statistics.

In the following tutorial, we confine ourselves to the *LM* test for *uno*rdered covariates (denoted LMuo), the *max*imally-selected *LM* test across the levels of an *or*dered covariate (denoted maxLMo), and the *max*imally-selected *LM* test for a continuous covariate (denoted maxLM). Another available option is the *d*ouble-*m*aximum test statistic (denoted DM), which is given by the maximum (over all respondents) of the maximum (over all item parameters) of the empirical cumulative process of the standardized individual score contributions (Merkle et al., [2013](#), [2014](#)). The double-maximum statistic is designed for detecting DIF effects with regard to continuous covariates. More details regarding the tests and their construction are provided by Merkle and Zeileis ([2013](#)) and Merkle et al. ([2014](#)), who also point out that other types of test statistics – e.g., (weighted) double maximum tests or Cramér-von-Mises-type tests – can be easily computed. Wang et al. ([2014](#)) give a practical introduction.

All of these statistics can be used to summarize the information of the invariance of various item parameters. The user can decide to summarize this information for the complete item set, leading to a DIF test for all items. Alternatively, this information can be summarized for individual items or groups of items, leading to item-wise DIF tests and DIF tests for groups of items, respectively.

The implementation of score-based measurement invariance tests within R

All functions needed for a DIF analysis with score-based tests are made available via three R packages, which are *psychotools*, *mirt* and *strucchange*. In this section, we will summarize their interaction and illustrate their application with example datasets.

The interaction of the R packages

The application of score-based measurement invariance tests conceptually consists of three major steps, which are: 1. the selection and estimation of a suitable IRT model, 2. the calculation of the individual score contributions, 3. the selection of a suitable test statistic and the calculation of the corresponding *p*-value. These steps correspond to specific steps in the analysis (see again Fig. 2):

- First, an appropriate IRT model has to be chosen based on theoretical considerations and on the type of items. If a Rasch-type model is chosen, for instance, the user can choose the CML or MML method for estimating the item parameters. If the model is specified correctly, both approaches lead to comparable results. On a technical level, *psychotools* allows the estimation of the item parameters with the CML method, whereas *mirt* allows the estimation of the item parameters using the MML approach.
- Second, the scores and individual score contributions are calculated. This is typically not done explicitly by the user, but done implicitly during the analysis. Like in the first step, these score contributions are calculated by *psychotools* using the CML approach and by *mirt* using the MML approach.
- Third, the user has to select a test statistic. This step is necessary for the calculation of a *p*-value via *strucchange*.

In summary, *mirt* and *psychotools* are used for the item parameter estimation and the calculation of the corresponding scores. *psychotools* is used for these calculations in the CML approach, whereas *mirt* is used with the MML approach. Based on these results, *strucchange* can be used for the calculation of *p*-values.

Important R functions

This section summarizes the most important functions of each R package on a conceptual level. We will present the functions in the order of the general workflow. Therefore, we will first address functions for the estimation of the item parameters, then functions for the score estimation, and functions for carrying out score-based measurement

invariance tests. Finally, we highlight some additional functions of the *psychotools* package. For all functions we also briefly introduce important arguments of the functions. To inspect the full list of arguments, the respective R help pages should be consulted. We start with some basic functions for fitting IRT models:

- `raschmodel`, `rsmodel`, `pcmodel`: These are basic fitting functions of the *psychotools* package for the binary Rasch model (`rasch`, Rasch, 1960) and two polynomial Rasch-type models, namely the rating scale model (`rsmodel`, Andrich, 1978) and the partial credit model (`pcmodel`, Masters, 1982). All of these functions estimate the item parameters within the CML approach. These functions require at least the response matrix provided via the argument `y`.
- `plmodel`, `gpcmodel`: These functions are user-friendly wrapper functions in *psychotools* that allow the estimation of common IRT models within the MML approach relying on *mirt*. `plmodel` allows the estimation of a wide range of models for dichotomous items, like the 2PL, 3PL (Birnbaum, 1968) and 4PL (Barton & Lord, 1981) models. `gpcmodel` allows the estimation of the generalized partial credit model (GPCM; Muraki, 1992). Both `plmodel` and `gpcmodel` also allow for fitting Rasch-type models within the MML approach, i.e., restricting the slopes to be equal for all items. Both functions require the response matrix as argument `y`. Modeling group differences in the true person parameter distributions, which is also named impact (e.g. Wang et al., 2012), is possible by the optional `impact` argument. This argument accepts a vector that defines a group membership for each respondent and models impact effects between the resulting groups. `plmodel` additionally accepts a `type` argument, which defines the estimated IRT model (e.g., by default `type = "2PL"`).
- *psychotools*' extractor functions: One advantage of using the fitting functions of the *psychotools* package, including the wrapper functions, is that *psychotools* provides plenty of easy to use extractor functions for all kinds of IRT parameters, i.e., item difficulty parameters, item discrimination parameters, lower and upper asymptote parameters, threshold parameters and also person parameters (`itempar`, `discrpar`, `guesspar`, `upperpar`, `threshpar` and `personpar`). Moreover, *psychotools* provides unified inference and visualization tools for its supported models.
- `sctest`: This function is part of the *strucchange* package. Based on the individual score contributions, a person covariate and the selection of a test statistic (or functional), this function allows for the calculation

of the test statistic and the associated *p*-value. This function is a generic function and in the scenarios described here, dispatches the `sctest.default` method, which has three main arguments: `x` is an estimated IRT model object, that is the return value of `mirt`, `raschmodel` or similar commands. `order.by` denotes the person covariate that is tested for DIF. `functional` allows the definition of the test statistic used for the score-based DIF test. Possible values for `functional` are, for instance, `"DM"` for the *double-maximum* statistic or `"LMuo"` for the *uno*rdered *L*agrange *M*ultiplier statistic, the *max*imally-selected *LM* test across the levels of an *o*rdered covariate (denoted `"maxLMo"`), or the *max*imally-selected *LM* test for a continuous covariate (denoted `"maxLM"`).

So far, we discussed some wrapper functions for fitting IRT models that are designed to be easy to use. The underlying *mirt* package provides the following, more flexible functions:

- `mirt`, `multipleGroup`: These functions are basic fitting functions for a wide range of unidimensional and multidimensional IRT models for dichotomous and polytomous items. The item parameter estimation is based on the MML approach. In the following, we assume that the EM algorithm (Bock & Aitkin, 1981; Baker & Kim, 2004) was employed to estimate the item parameters within the MML approach. When `mirt` is used for estimating the item parameters, the person parameter distribution is assumed to be identical for all subgroups of the population of respondents, i.e., the standard normal distribution. `multipleGroup` implements the multiple group estimation approach of Bock and Zimowski (1997) to allow the modeling of different person parameter distributions in various predefined groups of respondents. `mirt` requires at least three arguments. The first argument, `data`, is the response matrix. The second argument, `model`, typically indicates the number of latent dimensions, which is 1 for unidimensional models. The third argument, `itemtype`, indicates the type of estimated IRT model and defaults to `"2PL"`. Other options are, for instance, `"Rasch"`, `"3PL"` or `"4PL"`. `multipleGroup` additionally uses a `group` argument, which accepts a vector that defines a group membership for each respondent and allows to model impact effects between the different groups.

The following important function is used in the background:

- `estfun`: Methods for this generic function (which is part of the *sandwich* package, see Zeileis 2004, 2006) are contained in *mirt* as well as *psychotools* and allow for the calculation of the individual score

contributions. Here, their single argument is a model object corresponding to an estimated IRT model, e.g., the return value of `mirt` or `raschmodel`. Given the importance of the scores for the underlying statistical tests, they can be seen as a central building block of our software implementation. However, it is typically not necessary to use these functions explicitly during an analysis, as will be shown in the following tutorials.

Table 2 summarizes the main R functions needed to conduct score-based tests of measurement invariance with their most important arguments.

Illustrations with empirical data

In our first example, we investigate DIF for a Rasch model (Rasch, 1960) with regard to a continuous and a categorical covariate, and in the second example, for a GPCM (Muraki, 1992) with regard to an ordinal and a categorical variable. More detailed simulation studies are reported in earlier studies on score-based measurement invariance tests for different other IRT models (Debelak & Strobl, 2019; Komboz et al., 2018; Strobl et al., 2011; Strobl et al., 2015; Wang et al., 2018).

Before we begin our analysis, we load the three R packages *mirt*, *psychotools* and *strucchange* that are used in these examples. Before the first usage of these packages, it is necessary to install them. This can be done by using `install.packages(c("mirt", "psychotools", "strucchange"))`. After their installation, these packages can be loaded via:

```
> library("mirt")
> library("psychotools")
> library("strucchange")
```

After these packages have been loaded, we are ready to start our DIF analysis.

Investigating DIF for a continuous or categorical covariate for the rasch model: the verbal aggression dataset

The dataset that we are going to analyze in this subsection is the Verbal Aggression dataset. These data were thoroughly discussed in the book of De Boeck and Wilson (2004) and in the paper of De Boeck et al. (2011). In these studies, Rasch models were used to model this dataset, with item parameters estimated via MML. Another reason for selecting this dataset is that it is freely available in the *psychotools* package. We therefore load this dataset first:

```
> data("VerbalAggression",
+ package = "psychotools")
```

Table 2 Overview of main R functions to conduct score-based tests

Function (R package)	Purpose	Central arguments	Return value
Model Fitting			
raschmodel (*psychotools*)	fitting the binary Rasch model (CML)	y (response matrix)	object of class "raschmodel"
rsmodel (*psychotools*)	fitting the rating scale model (CML)	y (response matrix)	object of class "rsmodel"
pcmodel (*psychotools*)	fitting the partial credit model (CML)	y (response matrix)	object of class "pcmodel"
plmodel (*psychotools*)	fitting the 1PL, 2PL, 3PL, 3PLu, or 4PL model (MML, *mirt* wrapper)	y (response matrix), type (type of IRT model)	object of class "plmodel"
gpcmodel (*psychotools*)	fitting the GPCM (MML, *mirt* wrapper)	y (response matrix)	object of class "gpcmodel"
mirt (*mirt*)	fitting various IRT models (MML)	data (response matrix), model (either indicating the number of latent dimensions or loadings, constraints etc. of the model), itemtype (type of IRT model, can be specified itemwise)	object of class "SingleGroupClass"
multipleGroup (*mirt*)	fitting various multiple group IRT models (MML)	as in mirt above and additionally group (defines a group membership for each respondent)	object of class "MultipleGroupClass"
Calculation of *p*-values			
sctest (*strucchange*)	generic function for calculating a test statistic and the associated *p*-value. Dispatches sctest.default by default	x (model object), order.by (typically a vector of person covariates), functional (typically a string specifying the test statistic, see main text for examples)	object of class "htest"
Background function			
estfun (*sandwich*)	generic function for calculating the individual score contributions (methods are in *mirt* and *psychotools*)	x (model object)	matrix of the individual score contributions

This dataset contains responses from a questionnaire assessing verbal aggressive behavior. There are overall 24 items, which were worked on by 316 respondents (243 female). For details on this questionnaire, we refer to De Boeck and Wilson (2004) or De Boeck et al. (2011). Like De Boeck et al. (2011), we code the response "no" as 0 and all other responses as 1. Person covariates that are part of this dataset are gender and a score for anger. In the context of this paper, we are interested in testing for DIF with respect to the anger and gender variables.

VerbalAggression$resp2 contains the responses of all 314 respondents (as rows) to the 24 items (as columns), coded by 0 and 1, VerbalAggression$anger contains the anger scores for all respondents. VerbalAggression$gender provides information on the respondents' gender, using the categories female and male. We start our analysis by selecting an IRT model and an estimation approach for our analysis. Like De Boeck and Wilson (2004) and De Boeck et al. (2011), we base our analysis on the Rasch model for

dichotomous data, but focus on the first six items, which correspond to the responses to the first frustrating situation in the test ("A bus fails to stop for me."). These six items use three behaviors ("cursing", "scolding", "shouting") in the "doing" and "wanting" mode as a possible reaction to this situation. As an estimation approach, we choose CML estimation, since this method makes less strict assumptions about the distribution of the person parameters. We therefore use the *psychotools* package for estimating the item parameters. The following R command fits the Rasch model and stores the results in the raschmodel object RM_VA:

```
> RM_VA <- raschmodel(
+ VerbalAggression$resp2[, 1:6])
```

The following command carries out a score-based test with regard to the covariate anger. The order.by argument contains the covariate that should be used for DIF testing. As a test statistic, we use the maximum LM test, which corresponds to the argument maxLM:

```
> sctest(RM_VA, order.by =
```

```
+ VerbalAggression$anger,
+ functional = "maxLM")

M-fluctuation test

data: RM_VA f(efp) = 14.019, p-value =
0.2308
```

`sctest` returns the value of the test statistic as well as the associated *p*-value. As the result shows, there is no significant DIF effect with regard to anger.

We also apply a score-based test which tests the null hypothesis of measurement invariance with regard to gender. In the following, we set the argument `functional` to `LMuo` (i.e., we use the unordered LM test) to account for the nominal scale of our person covariate gender. This leads to:

```
> sctest(RM_VA,
+ order.by = VerbalAggression$gender,
+ functional = "LMuo")
M-fluctuation test

data: RM_VA f(efp) = 9.8788, p-value =
0.07874
```

This result does not indicate a significant DIF effect either.

Investigating DIF for an ordinal or categorical covariate for the generalized partial credit model: the generic conspiracist beliefs dataset

Our second example dataset is also available in the *psychotools* package and contains the responses of 2449 persons to 15 items measuring belief in conspiracy theories as well as two covariates, area (area one lived in as a child) and gender. The area covariate is of ordinal nature and takes on the values "rural", "suburban" and "urban".

As in the first example, we first select an IRT model for our analysis. Since our items are polytomous and our dataset is large enough to allow a sufficiently accurate estimation of slope parameters, we choose the GPCM. As this model does not belong to the family of Rasch-type models, we rely on MML estimation for estimating the parameters. A third decision concerns the use of a wrapper function for fitting the model. This decision does not affect the results of the analysis, but the wrapper functions included in the *psychotools* package are generally more convenient in their application. First, we load the dataset from the *psychotools* package:

```
> data("ConspiracistBeliefs2016",
+ package = "psychotools")
```

The responses are named `resp` and the variable for which impact is modeled is `area`. Both are stored in the data frame `ConspiracistBeliefs2016`. The following code estimates the GPCM item parameters while accounting for a possible impact effect with regard to area; we remind the reader that an impact effect corresponds to a shift in the distribution of the ability parameters. As was already argued in the introduction, this modeling choice avoids a possibly increased Type I error rate of the following score-based test. The model fitting is done with the `gpcmodel` function, which is one of the wrapper functions. We start our analysis by fitting the GPCM, which is done using *mirt* in the background. The `impact` argument specifies the variable for which impact is modeled. The `maxit` argument determines the maximum number of EM iterations for this estimation procedure. The default value in *mirt* (and analogously in the wrappers in *psychotools*) is 500, but larger datasets, as the one analyzed here, typically require a larger number of iterations. If the number of iterations is too low and the estimation algorithm does not converge, a warning will be printed.

```
> GPCM_Area <- gpcmodel(
+ ConspiracistBeliefs2016$resp,
+ impact = ConspiracistBeliefs2016$area,
+ maxit = 1000)
```

Based on the fitted model, a score-based test with regard to the area covariate (employing the maximally-selected LM test, `maxLMo`, across the levels of the ordered covariate) can be carried out via:

```
> sctest(GPCM_Area,
+    order.by = ConspiracistBeliefs2016$area,
+    functional = "maxLMo")

M-fluctuation test

data: GPCM_Area f(efp) = 89.605, p-value =
0.2205
```

This test is not significant. Therefore, we do not find a significant DIF effect with regard to the area covariate. It should be noted that we also could have proceeded to only test for DIF in individual items. In our example that uses the GPCM whose item parameters were estimated by MML estimation, this can be achieved by identifying the respective item parameters and providing their indices via the `parm` argument of `sctest`. For instance, if we wanted to test for DIF in only the first item, and therefore in the first five item parameters (that is, one slope and four threshold parameter), this could be done via:

```
> sctest(GPCM_Area,
+    order.by = ConspiracistBeliefs2016$area,
+    functional = "maxLMo", pars = 1:5)

M-fluctuation test
```

```
data: GPCM_Area f(efp) = 89.605, p-value
= 0.2206
```

The result is again not significant, which indicates an absence of DIF in the first five item parameters, which corresponds to the parameters of the first item. On a statistical level, this corresponds to a DIF test that checks whether the item parameters of the first item are invariant with regard to the area covariate while assuming that the remaining item parameters are stable. Please note that the order of the item parameters in the `pars` argument depends on the estimated model. It is equal to the order in the output of the `estfun` function, which can be directly inspected by the user. In practical applications of these item-wise DIF tests, the assumption that the parameters of all other items are stable should be critically reflected by the user before interpreting the resulting *p*-value. In the DIF testing literature, items that are set equal between the groups are termed anchor items, and it is important to select these items in a sensible way.

At this point, we also want to point to the `anchortest` function available in the *psychotools* package, which allows users to conduct an itemwise Wald test for pre-defined groups with "built-in" anchor selection based on a variety of methods. The anchor selection methods available in this function are described in Kopf et al. (2015a, b).

For completeness, we also test for DIF regarding the covariate gender. First, we fit the GPCM modeling potential impact in gender:

```
> GPCM_Gender <- gpcmodel(
+ ConspiracistBeliefs2016$resp,
+ impact = ConspiracistBeliefs2016$gender,
+ maxit = 1000)
```

Second, we conduct the score-based test for all items with the unordered Lagrange multiplier statistic, using the functional `LMuo`:

```
> sctest(GPCM_Gender,
+ order.by =
+ ConspiracistBeliefs2016$gender,
+ functional = "LMuo")


M-fluctuation test

data: GPCM_Gender f(efp) = 288.61, p-value
= 7.993e-11
```

This test finds a *p*-value smaller than 0.001. We therefore reject the null hypothesis of measurement invariance with regard to gender. We conclude that at least one item parameter differs between females and males.

Discussion

In this tutorial, we provided a detailed introduction to the theory of score-based tests for detecting DIF in IRT models, as well as their application within the R system for statistical computing relying on the packages *psychotools*, *mirt* and *strucchange*. The underlying family of statistical tests was developed for investigating the invariance of ML estimators with regard to categorical, ordinal and continuous covariates. It can be applied to both CML and MML estimated IRT models.

New functions that we have added to the *psychotools* and *mirt* packages make these kind of score-based tests of measurement invariance accessible to the general user. Easy-to-use wrapper functions included in *psychotools* now allow for a straightforward application of these tests without requiring a detailed knowledge of the underlying R commands. So far, these wrapper functions are available for the Rasch model (with MML estimation; 1PL model), the 2PL and 3PL models of Birnbaum (1968), the more general 4PL model (Barton & Lord, 1981), as well as the GPCM (Muraki, 1992), including the partial credit model (Masters, 1982). Alternatively, users can use *mirt* directly for an application of these tests to all kinds of IRT models estimated within the MML approach. These applications might be particularly interesting for substantive researchers. Future extensions might contain additional wrapper functions, for instance for multidimensional models like the multidimensional compensatory 2PL model (Reckase, 2009). Since it is already possible to obtain MML estimates for these models with the *mirt* package, it is also possible to check their item parameter invariance with score-based measurement invariance tests, although no simulation results regarding the power of these tests have been published so far for these models. By similar means it is also possible to carry out a power analysis of score-based measurement invariance tests as part of the IRT analysis of empirical data. This application might be of interest for psychometricians who wish to investigate the psychometric characteristics of these tests. The *psychotools* package contains a vignette on how to conduct simulation studies with score-based tests.

We want to conclude this paper with a discussion of possible extensions of the presented method. In principle, these extensions could either address situations where assumptions underlying the methods presented here are violated, or they could extend the functionality of the software.

We start with extensions of the first type. This paper considers a scenario where all respondents worked 1) on the same item set and 2) a likelihood-based approach is used

for estimating the item parameters. This scenario is fairly general and covers a wide range of situations encountered in empirical data analysis. A first extension of practical interest is the development of analogous methods for Bayesian estimation methods, which are also commonly used in IRT analysis (Fox, 2010; Levy & Mislevy, 2016). A second extension could be a scenario where not all respondents worked on the same item set. A typical example is the application of a computerized adaptive test or a multistage test (Van der Linden & Glas, 2010; Yan et al., 2016). These techniques of item presentation are well-known in educational settings and have in common that the selection of items worked on during a test (partly) depends on previous test behavior, whereas the item parameters are typically treated as known.

We also mention some possible extensions to the functionality of the software. An example concerns the development of tree-based methods based on recursive model-based partitioning. This type of method is based on simultaneous tests for measurement invariance with regard to multiple person covariates. It aims at a data-driven detection of groups for which no significant DIF effects can be detected. For Rasch-type models and the CML approach, such methods are already provided in the *psychotree* package (Komboz et al., 2018; Strobl et al., 2015).

Another field for future work could be the inclusion and evaluation of sensitivity statistics. These statistics indicate the expected change of model parameters if they are no longer assumed to be invariant across groups of respondents. After score-based tests have detected a violation of invariance, statistics of this kind could help to investigate the practical relevance of this model violation. An example is the expected parameter change (EPC)-interest statistic, which was suggested by Oberski (2014) for structural equation models. A closely related idea is the inclusion of effect sizes for DIF effects (Chalmers, 2018; Steinberg & Thissen, 2006).

Computational Details All results were obtained using the R system for statistical computing (R Core Team, 2019) version 3.5.3, running under Windows 10 x64 (build 18363) and employing the add-on packages *mirt* (Chalmers, 2012) version 1.31, *psychotools* (Zeileis et al., 2021) version 0.7-0, *sandwich* (Zeileis, 2006) version 3.0-0, and *strucchange* (Zeileis et al., 2002) version 1.5-2, which are freely available under the General Public License from the Comprehensive R Archive Network at https://cran.r-project.org/. Numerical values were rounded based on the IEC 60559 standard.

Open Practices Statement The datasets used in this study are freely available in the add-on package *psychotools* (Zeileis et al., 2021).

References

American Educational Research Association. (2014). *Standards for educational and psychological testing*. Washington: American Educational Research Association.

Andrews, D. W. K. (1993). Tests for parameter instability and structural change with unknown change point. *Econometrica*, *61*(4), 821–856. https://doi.org/10.2307/2951764

Andrich, D. (1978). Application of a psychometric rating model to ordered categories which are scored with successive integers. *Applied Psychological Measurement*, *2*(4), 581–594. https://doi.org/10.1177/014662167800200413

Arnold, M., Oberski, D. L., Brandmaier, A. M., & Voelkle, M. C. (2020). Identifying heterogeneity in dynamic panel models with individual parameter contribution regression. *Structural Equation Modeling: A Multidisciplinary Journal*, *27*(4), 613–628. https://doi.org/10.1080/10705511.2019.1667240

Baker, F. B., & Kim, S.-H. (2004). *Item response theory: Parameter estimation techniques*, (2nd ed.). Boca Raton: CRC Press.

Barton, M. A., & Lord, F. M. (1981). An upper asymptote for the three-parameter logistic item-response model. *ETS Research Report Series, 1981* (1), i–8. https://doi.org/10.1002/j.2333-8504.1981.tb01255.x.

Bauer, D. J. (2017). A more general model for testing measurement invariance and differential item functioning. *Psychological Methods*, *22*(3), 507–526. https://doi.org/10.1037/met0000077

Bauer, D. J., & Hussong, A. M. (2009). Psychometric approaches for developing commensurate measures across independent studies: Traditional and new models. *Psychological Methods*, *14*(2), 101–125. https://doi.org/10.1037/a0015583

Billingsley, P. (1995). *Probability and measure*. New York: Wiley.

Birnbaum, A. (1968). Some latent trait models and their use in inferring an examinee's ability. In Lord, F. M., & Novick, M. R. (Eds.) *Statistical theories of mental test scores*. Boston: Addison-Wesley.

Bock, R. D. (1972). Estimating item parameters and latent ability when responses are scored in two or more nominal categories. *Psychometrika*, *37*(1), 29–51. https://doi.org/10.1007/bf02291411

Bock, R. D., & Aitkin, M. (1981). Marginal maximum likelihood estimation of item parameters: Application of an EM algorithm. *Psychometrika*, *46*(4), 443–459. https://doi.org/10.1007/bf02293801

Bock, R. D., & Zimowski, M. F. (1997). Multiple group IRT. In Van der Linden, W. J., & Hambleton, R. K. (Eds.) *Handbook of modern item response theory*, (pp. 433–448). New York: Springer.

Borsboom, D. (2006). The attack of the psychometricians. *Psychometrika*, *71*(3), 425–440. https://doi.org/10.1007/s11336-006-1447-6

Brotherton, R., French, C. C., & Pickering, A. D. (2013). Measuring belief in conspiracy theories: The Generic Conspiracist Beliefs scale. *Frontiers in Psychology*, *4*, 279. https://doi.org/10.3389/fpsyg.2013.00279

Camilli, G. (2006). Test fairness. In Brennan, R. (Ed.) *Educational measurement*. (4th ed., pp. 221–256). Westport: American Council on Education and Praeger.

Chalmers, R. P. (2012). mirt: A multidimensional item response theory package for the R environment. *Journal of Statistical Software, 48* (6), 1–29. https://doi.org/10.18637/jss.v048.i06.

Chalmers, R. P. (2018). Model-based measures for detecting and quantifying response bias. *Psychometrika, 83*(3), 696–732. https://doi.org/10.1007/s11336-018-9626-9

Dai, Y. (2013). A mixture Rasch model with a covariate: A simulation study via Bayesian Markov Chain Monte Carlo estimation. *Applied Psychological Measurement, 37*(5), 375–396. https://doi.org/10.1177/0146621612475076

Debelak, R., & Strobl, C. (2019). Investigating measurement invariance by means of parameter instability tests for 2PL and 3PL models. *Educational and Psychological Measurement, 79*(2), 385–398. https://doi.org/10.1177/0013164418777784

De Boeck, P., Bakker, M., Zwitser, R., Nivard, M., Hofman, A., Tuerlinckx, F., & Partchev, I. (2011). The estimation of item response models with the lmer function from the lme4 package in R. *Journal of Statistical Software, 39*(12), 1–28. https://doi.org/10.18637/jss.v039.i12

De Boeck, P., & Wilson, M. (2004). *Explanatory item response models*. New York,: Springer.

Falk, C. F., & Cai, L. (2016). Maximum marginal likelihood estimation of a monotonic polynomial generalized partial credit model with applications to multiple group analysis. *Psychometrika, 81*(2), 434–460. https://doi.org/10.1007/s11336-014-9428-7

Fox, J.-P. (2010). *Bayesian item response modeling: Theory and applications*. New York: Springer.

Glas, C. A. W. (1998). Detection of differential item functioning using Lagrange multiplier tests. *Statistica Sinica, 8*(3), 647–667.

Glas, C. A. W. (1999). Modification indices for the 2-PL and the nominal response model. *Psychometrika, 64*(3), 273–294. https://doi.org/10.1007/BF02294296

Glas, C. A. W., & Suárez Falcón, J. C. (2003). A comparison of item-fit statistics for the three-parameter logistic model. *Applied Psychological Measurement, 27*(2), 87–106. https://doi.org/10.1177/0146621602250530

Holland, P. W., & Wainer, H. (1993). *Differential item functioning*. New York: Taylor & Francis.

Komboz, B., Strobl, C., & Zeileis, A. (2018). Tree-based global model tests for polytomous Rasch models. *Educational and Psychological Measurement, 78*(1), 128–166. https://doi.org/10.1177/0013164416664394

Kopf, J., Zeileis, A., & Strobl, C. (2015a). Anchor selection strategies for DIF analysis: Review, assessment, and new approaches. *Educational and Psychological Measurement, 75*(1), 22–56. https://doi.org/10.1177/0013164414529792

Kopf, J., Zeileis, A., & Strobl, C. (2015b). A framework for anchor methods and an iterative forward approach for DIF detection. *Applied Psychological Measurement, 39*(2), 83–103. https://doi.org/10.1177/0146621614544195

Levy, R., & Mislevy, R. J. (2016). *Bayesian psychometric modeling*. Boca Raton: Chapman and Hall/CRC.

Li, T., Jiao, H., & Macready, G. B. (2016). Different approaches to covariate inclusion in the mixture Rasch model. *Educational and Psychological Measurement, 76*(5), 848–872. https://doi.org/10.1177/0013164415610380

Liu, Y., Magnus, B. E., & Thissen, D. (2016). Modeling and testing differential item functioning in unidimensional binary item response models with a single continuous covariate: A functional data analysis approach. *Psychometrika, 81*(2), 371–398. https://doi.org/10.1007/s11336-015-9473-x

Magis, D., Tuerlinckx, F., & De Boeck, P. (2010). A general framework and an R package for the detection of dichotomous differential item functioning. *Behavior Research Methods, 42*(3), 847–862. https://doi.org/10.3758/BRM.42.3.847

Masters, G. N. (1982). A Rasch model for partial credit scoring. *Psychometrika, 47*(2), 149–174. https://doi.org/10.1007/BF02296272

Maydeu-Olivares, A., Hernández, A., & McDonald, R. P. (2006). A multidimensional ideal point item response theory model for binary data. *Multivariate Behavioral Research, 41* (4), 445–472. https://doi.org/10.1207/s15327906mbr4104n 2.

Merkle, E. C., Fan, J., & Zeileis, A. (2014). Testing for measurement invariance with respect to an ordinal variable. *Psychometrika, 79*(4), 569–584. https://doi.org/10.1007/s11336-013-9376-7

Merkle, E. C., & Zeileis, A. (2013). Tests of measurement invariance without subgroups: A generalization of classical methods. *Psychometrika, 78*(1), 59–82. https://doi.org/10.1007/s11336-012-9302-4

Millsap, R. E. (2011). *Statistical approaches to measurement invariance*. Routledge: New York.

Molenaar, D. (2020). A exible moderated factor analysis approach to test for measurement invariance across a continuous variable. *Psychological Methods*. https://doi.org/10.1037/met0000360.

Molenaar, I. W. (1995). Estimation of item parameters. In Fischer, G. H., & Molenaar, I. W. (Eds.) *Rasch models: Foundations, recent developments, and applications*, (pp. 39–51). New York: Springer.

Moustaki, I. (2003). A general class of latent variable models for ordinal manifest variables with covariate effects on the manifest and latent variables. *British Journal of Mathematical and Statistical Psychology, 56*(2), 337–357. https://doi.org/10.1348/000711003770480075

Muraki, E. (1992). A generalized partial credit model: Application of an EM algorithm. *Applied Psychological Measurement, 16*(2), 159–176. https://doi.org/10.1002/j.2333-8504.1992.tb01436.x

Oberski, D. L. (2013). A exible method to explain differences in structural equation model parameters over subgroups. Retrieved from http://daob.nl/wp-content/uploads/2013/06/SEM-IPC-manuscript-new.pdf.

Oberski, D. L. (2014). Evaluating sensitivity of parameters of interest to measurement invariance in latent variable models. *Political Analysis, 22*(1), 45–60. https://doi.org/10.1093/pan/mpt014

Open Source Psychometrics Project (2016). Data from: The Generic Conspiracist Beliefs scale [Dataset]. Retrieved from https://openpsychometrics.org/_rawdata/.

R Core Team (2019). R: A language and environment for statistical computing [Computer software manual]. Vienna, Austria. Retrieved from https://www.R-project.org/.

Rasch, G. (1960). *Probabilistic models for some intelligence and achievement tests*. Copenhagen: Danish Institute for Educational Research.

Reckase, M. (2009). *Multidimensional item response theory*. New York: Springer.

Roberts, J. S., Donoghue, J. R., & Laughlin, J. E. (2000). A general item response theory model for unfolding unidimensional polytomous responses. *Applied Psychological Measurement, 24*(1), 3–32. https://doi.org/10.1177/01466216000241001

Rost, J. (1990). Rasch models in latent classes: an integration of two approaches to item analysis. *Applied Psychological Measurement, 14*(3), 271–282. https://doi.org/10.1177/014662169001400305

Samejima, F. (1969). Estimation of latent ability using a response pattern of graded scores. *Psychometrika, 34*(1 Supplement), 1–97. https://doi.org/10.1007/BF03372160

Steinberg, L., & Thissen, D. (2006). Using effect sizes for research reporting: examples using item response theory to analyze differential item functioning. *Psychological Methods, 11*(4), 402–415. https://doi.org/10.1037/1082-989X.11.4.402

Strobl, C., Kopf, J., & Zeileis, A. (2015). Rasch trees: A new method for detecting differential item functioning in the Rasch model. *Psychometrika*, *80*(2), 289–316. https://doi.org/10.1007/s11336-013-9388-3

Strobl, C., Wickelmaier, F., & Zeileis, A. (2011). Accounting for individual differences in Bradley-Terry models by means of recursive partitioning. *Journal of Educational and Behavioral Statistics*, *36*(2), 135–153. https://doi.org/10.3102/1076998609359791

Tay, L., Newman, D. A., & Vermunt, J. K. (2011). Using mixed-measurement item response theory with covariates (MM-IRT-C) to ascertain observed and unobserved measurement equivalence. *Organizational Research Methods*, *14*(1), 147–176. https://doi.org/10.1177/1094428110366037

Van der Linden, W. J., & Glas, C. A. W. (2010). *Elements of adaptive testing*. New York: Springer.

Wang, T., Merkle, E. C., & Zeileis, A. (2014). Score-based tests of measurement invariance: Use in practice. *Frontiers in Psychology*, *5*, 438. https://doi.org/10.3389/fpsyg.2014.00438

Wang, T., Strobl, C., Zeileis, A., & Merkle, E. C. (2018). Score-based tests of differential item functioning via pairwise maximum likelihood estimation. *Psychometrika*, *83*(1), 132–155. https://doi.org/10.1007/s11336-017-9591-8

Wang, W.-C., Shih, C.-L., & Sun, G.-W. (2012). The DIF-free-then-DIF strategy for the assessment of differential item functioning. *Educational and Psychological Measurement*, *72*(4), 687–708. https://doi.org/10.1177/0013164411426157

Yan, D., Von Davier, A. A., & Lewis, C. (2016). *Computerized multistage testing: Theory and applications*. Boca Raton: CRC Press.

Zeileis, A. (2004). Econometric computing with HC and HAC covariance matrix estimators. *Journal of Statistical Software*, *11*(10), 1–17. https://doi.org/10.18637/jss.v011.i10

Zeileis, A. (2006). Object-oriented computation of sandwich estimators. *Journal of Statistical Software*, *16*(9), 1–16. https://doi.org/10.18637/jss.v016.i09

Zeileis, A., & Hornik, K. (2007). Generalized M-fluctuation tests for parameter instability. *Statistica Neerlandica*, *61*(4), 488–508. https://doi.org/10.1111/j.1467-9574.2007.00371.x

Zeileis, A., Leisch, F., Hornik, K., & Kleiber, C. (2002). strucchange: An R package for testing for structural change in linear regression models. *Journal of Statistical Software, 7* (1), 1–38. https://doi.org/10.18637/jss.v007.i02

Zeileis, A., Strobl, C., Wickelmaier, F., Komboz, B., Kopf, J., Schneider, L., & Debelak, R. (2021). psychotools: Infrastructure for psychometric modeling [Computer software manual]. Retrieved from https://CRAN.R-project.org/package=psychotools (R package version 0.7-0).